

Process vs. Outcome Reward: Which is Better for Agentic RAG Reinforcement Learning

Wenlin Zhang^{1*}, Xiangyang Li^{2*}, Kuicai Dong^{3*}, Yichao Wang², Pengyue Jia¹,
Xiaopeng Li¹, Yingyi Zhang¹, Derong Xu¹, Zhaocheng Du², Huifeng Guo²,
Ruiming Tang², Xiangyu Zhao¹

¹ City University of Hong Kong, Hong Kong, ² Noah's Ark Lab, Huawei, China

³ Noah's Ark Lab, Huawei, Singapore

{w.l.z, jia.pengyue, xiaopli2-c, yzhang6375-c, derongxu2-c}@my.cityu.edu.hk,

{lixiangyang34, dong.kuicai, wangyichao5, zhaochengdu, huifeng.guo, tangruiming}@huawei.com,

xianzhao@cityu.edu.hk

Abstract

Retrieval-augmented generation (RAG) enhances the text generation capabilities of large language models (LLMs) by integrating external knowledge and up-to-date information. However, traditional RAG systems are limited by static workflows and lack the adaptability required for multistep reasoning and complex task management. To address these limitations, agentic RAG systems (e.g., DeepResearch) have been proposed, enabling dynamic retrieval strategies, iterative context refinement, and adaptive workflows for handling complex search queries beyond the capabilities of conventional RAG. Recent advances, such as Search-R1, have demonstrated promising gains using outcome-based reinforcement learning, where the correctness of the final answer serves as the reward signal. Nevertheless, such outcome-supervised agentic RAG methods face challenges including low exploration efficiency, gradient conflict, and sparse reward signals. To overcome these challenges, we propose to utilize fine-grained, process-level rewards to improve training stability, reduce computational costs, and enhance efficiency. Specifically, we introduce a novel method ReasonRAG that automatically constructs RAG-ProGuide, a high-quality dataset providing process-level rewards for (i) query generation, (ii) evidence extraction, and (iii) answer generation, thereby enhancing model inherent capabilities via process-supervised reinforcement learning. With the process-level policy optimization, the proposed framework empowers LLMs to autonomously invoke search, generate queries, extract relevant evidence, and produce final answers. Compared to existing approaches such as Search-R1 and traditional RAG systems, ReasonRAG, leveraging RAG-ProGuide, achieves superior performance on five benchmark datasets using only 5k training instances, significantly fewer than the 90k training instances required by Search-R1. Our code is available at <https://github.com/wlzhang2020/ReasonRAG>.

1 Introduction

Large language models (LLMs) [1, 2, 3] have demonstrated substantial proficiency in text generation and natural language understanding tasks [4, 5]. However, their reliance on static training data constrains their ability to address dynamic and real-time queries, often resulting in outdated or hallucinated information [6, 7, 8]. Retrieval-Augmented Generation (RAG) [9] has emerged as a promising solution by equipping LLMs with external knowledge sources, improving the relevance,

*Equal contribution




Figure 1: Outcome-supervised vs. process-supervised RL for multi-step reasoning. Each circle denotes one-step reasoning response, where the correct response is colored as "green" and error response is colored as "red".

factual accuracy, and timeliness of responses [10, 11, 12]. Despite these advances, traditional RAG architectures are limited by their linear and static workflows, which suffer from complex multi-step reasoning, deep contextual integration, and iterative response refinement [9, 13]. To address these shortcomings, agentic RAG (*e.g.*, DeepResearch [14, 15]) systems have been developed, enabling dynamic retrieval strategies, enhanced contextual understanding, and iterative refinement. Achieving agentic RAG requires the underlying LLMs to orchestrate retrieval, filter relevant information, and iteratively refine their outputs, resulting in more adaptive and efficient information processing.

To advance agentic RAG, early approaches [16, 17, 18, 19, 20] primarily focused on prompt engineering to adapt powerful LLMs to agentic workflows. However, due to the limited reasoning and instruction-following capabilities of LLMs, supervised fine-tuning (SFT) methods [21] have been introduced, extending prompt-based approaches by directly optimizing and refining model parameters. Due to SFT storing reasoning steps within the model parameters, the improved reasoning capabilities often encounter challenges in generalizing across different domains [22]. More recently, reinforcement learning (RL) methods (*e.g.*, OpenAI-O1 [23] and DeepSeek-R1 [24]) achieve notable improvements in LLM reasoning by employing outcome-supervised RL techniques. Building on these developments, Search-R1 [25] incorporates a search engine as part of the LLM’s environment and leverages outcome-based reinforcement learning, using the correctness of the final answer as the reward signal. These advances demonstrate that outcome-supervised reinforcement learning can substantially enhance the capabilities required for agentic RAG, enabling straightforward, end-to-end optimization of the entire workflow.

Despite its promise, outcome-supervised RL also presents inherent limitations, as illustrated in Figure 1. First, **low exploration efficiency** occurs since the model must generate a complete reasoning chain before receiving any reward [25]. Ideally, the reward should be given when errors occur at intermediate steps to facilitate learning. Second, **gradient conflict** arises when mistakes occur late in the reasoning process; the entire sequence (including correct early steps) is penalized [26]. This can lead to conflicting gradients that can push correct actions in the wrong direction. Third, **the rewards are sparse**, as outcome-supervised RL only provides feedback upon producing the final answer [27]. Reward sparsity relies on more training data and steps to converge, as the model receives infrequent learning signals. In contrast, **process-supervised RL** addresses these issues by providing fine-grained, stepwise rewards throughout the reasoning process, enabling more efficient exploration, reducing gradient conflict, and accelerating model learning through denser feedback.

However, applying process-supervised RL to RAG presents several key challenges. (1) **Process Reward Design**: Effective process rewards are essential for guiding the model toward the shortest and most efficient correct reasoning path. Rewards must incentivize helpful intermediate steps that lead to the correct final answer, while penalizing unnecessarily long or circuitous reasoning sequences [28]. (2) **Exploration Efficiency and Annotation Cost**: While human annotators who are skilled in information retrieval can create high-quality process-level annotations by decomposing complex retrieval tasks into efficient steps, this approach is prohibitively expensive due to the substantial manual effort involved [29]. In contrast, autonomous RAG agents can generate a wide range of possible retrieval and reasoning steps, but this large search space makes it difficult to identify and select high-quality, meaningful steps for use as process-level annotations.

To address these challenges, we propose ReasonRAG, an advanced process-supervised RL method to enhance agentic RAG reasoning. Specifically, ReasonRAG employs Monte Carlo Tree Search (MCTS) [30] as a search strategy to efficiently balance exploration and exploitation, enabling thorough exploration of diverse reasoning paths and identification of high-reward intermediate steps for guiding

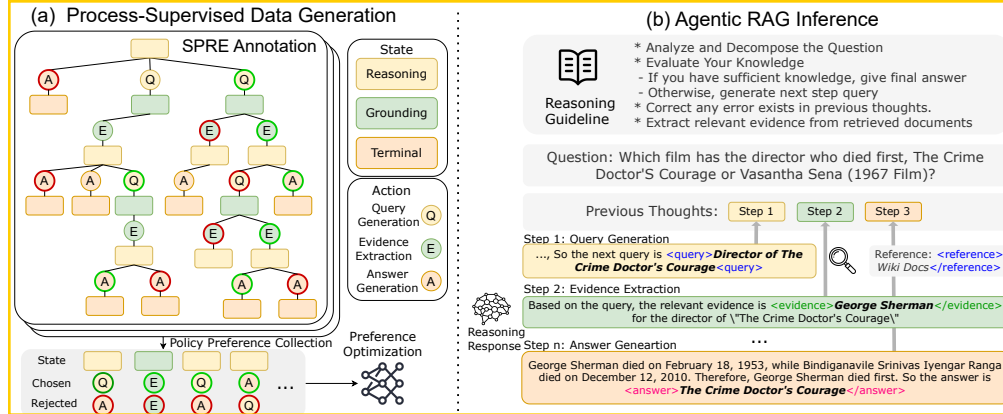


Figure 2: Framework of ReasonRAG. Figure (a) illustrates the policy optimization based on process supervision. MCTS guides the construction of the state-action tree and the assignment of process-level rewards for fine-grained policy optimization. (Actions derived from the same state are color-coded by reward: green circle (highest), red circle (lowest).) Figure (b) demonstrates an inference example.

the RAG process. Building on these paths, we introduce a novel **Shortest Path Reward Estimation (SPRE)** algorithm to assign rewards. SPRE favors sequences that lead to the correct answer while penalizing unnecessarily lengthy reasoning, thereby promoting efficiency. This approach yields **RAG-ProGuide**, a dataset comprising 5k queries with 13,000 high-quality process-level preference pairs. Using RAG-ProGuide and our process-supervised Direct Preference Optimization (DPO) [31] strategy, ReasonRAG is further trained to make dynamic decisions, such as whether to invoke retrieval, formulate subsequent search queries, analyze retrieved documents for relevant evidence, and synthesize evidence into final answers. Extensive experiments on five benchmark RAG datasets show ReasonRAG (trained with only 13k process-level steps) outperforms Search-R1 (trained on 90k queries with approximately 270k intermediate steps), suggesting the superiority of process-supervised RL over outcome-supervised RL. Our key contributions can be summarized as follows:

- We propose ReasonRAG, an automatic framework for agentic RAG process-level reward annotations. We introduce SPRE for efficient RAG process-level reward annotation and MCTS for high-quality decision space exploration.
- We introduce a process-level annotation dataset RAG-ProGuide, which serves as an off-policy dataset, and can be easily applied for various LLM policy optimization.
- We conduct extensive comparative experiments of outcome-supervised RL and process-supervised RL for RAG reasoning with Owen2.5-7B-Instruct. The experimental results on five benchmark datasets demonstrate the superiority and training efficiency of ReasonRAG.

2 ReasonRAG Framework

2.1 Framework Overview

This section details the design of ReasonRAG framework, as depicted in Figure 2. Figure 2a outlines our approach for constructing high-quality process-supervised data. We first introduce *Shortest Path Reward Estimation (SPRE)* to provide process-level supervision reward (see Section 2.2.1). To efficiently gather these rewards, we employ Monte Carlo Tree Search (MCTS) algorithm to explore the vast decision space in agentic RAG and collect informative intermediate steps (see Section 2.2.2).

The resulting process-supervised dataset, RAG-ProGuide (see Section 2.2.3), is then used to optimize ReasonRAG via *policy preference optimization*. This training strategy guides the model to prefer desirable reasoning trajectories in agentic RAG (see Section 2.3).

Figure 2b illustrates the agentic RAG inference workflow in ReasonRAG. During inference, the model adaptively conducts reasoning by dynamically invoking search engine and interleaving three core actions: query generation, evidence extraction, and answer generation (see Section 2.4).

2.2 Process-Supervised Data Generation

Effective process-supervised policy optimization requires high-quality, granular reward signals at the process level. As outlined in Section 1, generating such rewards for agentic RAG presents two main challenges: (1) the absence of reward functions for intermediate reasoning steps, and (2) the lack of an efficient and cost-effective method to generate informative reasoning trajectories. To overcome these challenges, we introduce a novel process-level reward function, SPRE (see Section 2.2.1), specifically designed for agentic RAG. Furthermore, we develop an MCTS-based approach (see Section 2.2.2) to efficiently explore the decision space and collect high-quality process-level data.

2.2.1 Shortest Path Reward Estimation (SPRE)

Unlike outcome-level rewards, process-level rewards provide supervision at each intermediate step of agentic RAG. A key challenge is the absence of ground-truth reward signals for partial reasoning trajectories. Furthermore, due to the large decision space, the reward function must account for both final correctness and reasoning efficiency. To address these challenges, we propose *Shortest Path Reward Estimation (SPRE)*, which evaluates the quality of each intermediate reasoning path by simulating its possible outcomes and penalizing unnecessarily long trajectories.

Formally, the agentic RAG process consists of an n -step sequence $[y_1, \dots, y_n]$, where each y_i represents the output of a single reasoning step, conditioned on the initial question x and previous steps $y_{<i}$. To evaluate a partial sequence $y_{1:t}$, we simulate multiple continuations, known as *rollouts*, until a final answer is obtained. By repeating the rollout process k times and scoring each outcome, we approximate the reward as a Monte Carlo-style estimation with step-based penalties:

$$Q_t = \text{MonteCarlo}(x, y_{1:t}) = \frac{1}{k} \sum_{i=1}^k v(\text{rollout}_i) \cdot \alpha^{\text{step}(\text{rollout}_i)} \quad (1)$$

Here, rollout_i is the i -th simulated completion of $y_{1:t}$, $v(\text{rollout}_i) \in [0, 1]$ denotes the correctness score (e.g., F1 match to the ground truth), and $\text{step}(\text{rollout}_i)$ is the number of total reasoning steps in the trajectory. The decay factor $\alpha \in (0, 1]$ penalizes unnecessarily long reasoning paths. This reward encourages the model to favor trajectories that achieve correct answers with fewer steps, thus balancing accuracy and efficiency in agentic RAG reasoning.

2.2.2 Monte Carlo Tree Search (MCTS) for Process-level Exploration

Although SPRE offers reliable reward signals for evaluating intermediate steps, generating diverse yet meaningful trajectories remains challenging. The search space in agentic RAG is extensive due to open-ended nature of retrieval, which requires continuous refinement of search queries for relevant information. To address this, we propose a tailored MCTS framework for agentic RAG. MCTS enables efficient exploration by selectively expanding the most promising reasoning paths based on estimated rewards.

We adapt MCTS to agentic RAG context by explicitly defining states and actions for tree construction. Formally, each intermediate reasoning step is represented as a state $s = (x, y_{<i}, \text{stage})$, where x is the original question, $y_{<i}$ is the sequence of prior reasoning outputs, and $\text{stage} \in \{\text{Reasoning}, \text{Grounding}, \text{Terminal}\}$ indicates current point of agentic flow. Actions $a \in \{\text{Query Generation}, \text{Evidence Extraction}, \text{Answer Generation}\}$ are determined by the current stage as follows:

- **Reasoning** stage: Choose between generating a new query for document retrieval or directly generating an answer. If a new query is generated, a retrieval operation is performed, and the retrieved documents are appended to the state for subsequent decisions. If an answer is produced, the process transitions to the *Terminal* stage.
- **Grounding** stage: Select evidence spans from the retrieved documents. Afterwards, the system returns to the *Reasoning* stage for further reasoning or answering.
- **Terminal** stage: End the exploration process when the final answer has been generated.

Based on the current state s , the policy for generating the next action a is defined as:

$$\pi(a | s) = \text{LLM}(a | s) = \begin{cases} \pi_\theta(\cdot | x, y_{<i}, p_{\text{stage}}), & \text{if stage is Reasoning} \\ \pi_\theta(\cdot | x, y_{<i}, \text{docs}, p_{\text{stage}}), & \text{otherwise} \end{cases} \quad (2)$$

Statistics	Number
Questions	4603
- PopQA	704 (15.3%)
- HotpotQA	2843 (61.8%)
- 2WikiMultihopQA	1056 (22.9%)
Actions	13289
- Query Generation	3295 (24.8%)
- Evidence Extraction	4305 (32.4%)
- Answer Generation	5689 (42.8%)
Avg./Min./Med./Max. Iteration	2.7/1/3/5
Avg./Min./Med./Max. Tokens	65.5/9/60/625

Table 1: Overall Dataset Statistics




Figure 3: Dataset Distribution.

State transitions are defined as $s_{t+1} = \text{concatenate}(s_t, a_t)$, where each action leads to a new node in the search tree, representing an extended reasoning sequence. This recursive process incrementally builds a tree rooted at the original question. MCTS then operates iteratively, performing three core steps: *selection*, *expansion*, and *backpropagation*. Specifically, at each iteration, MCTS selects promising paths using a Upper Confidence Bound (UCB) based objective, expands new states by sampling LLM-generated actions, and backpropagates SPRE-estimated rewards (see Equation 1) to update the tree (see more comprehensive explanation of MCTS exploration in Appendix A). Integrating MCTS with SPRE enables efficient exploration and prioritization of high-reward reasoning steps, producing high-quality process-level annotations to optimize the agentic RAG policy.

2.2.3 RAG-ProGuide Dataset

Construction. Using the MCTS-based exploration framework, we construct a high-quality process-supervised dataset RAG-ProGuide to facilitate process-level policy optimization. We randomly sample 3,000 questions each from PopQA [32], HotpotQA [33], and 2WikiMultihopQA [34], covering both single-hop and multi-hop question answering tasks. GPT-4o [35] serves as the policy model to simulate the agentic RAG reasoning process within the MCTS framework (see Section 2.2.2). During tree search, we prune all branches that do not yield a final answer. For each complete trajectory, we compute the F_1 score between the predicted answer and ground truth, and use this correctness signal to estimate intermediate node rewards via SPRE (see Section 2.2.1). These rewards are propagated through the MCTS tree to guide preference pair selection. To ensure high-quality preference data, we perform post-processing to remove duplicates and uninformative comparisons: (i) we discard preference pairs with identical response sequences, and (ii) pairs with a reward difference less than 0.01. After filtering, the final dataset consists of 4,603 questions and 13,289 distinct preference pairs.

Dataset Statistics. Table 1 presents detailed statistics and distributions for our dataset. The questions are drawn from PopQA, HotpotQA, and 2WikiMultihopQA, providing comprehensive coverage of both single-hop and multi-hop reasoning within the RAG decision space. The dataset contains a balanced distribution of three reasoning actions, reflecting the multi-stage nature of the agentic RAG process. As shown in Figure 2a, the distribution of preference pair types demonstrates diverse comparative scenarios: the x-axis abbreviations (A: answer generation, Q: query generation, E: evidence extraction) indicate action types in accepted versus rejected paths. This diversity ensures fine-grained comparative coverage across different reasoning stages. Figure 2b indicates a range of reasoning iteration counts, consistent with the complexity of multi-hop inference. Figure 2c shows a broad distribution of response token lengths, confirming the dataset’s capacity to capture various response complexities. Additionally, Figure 2d highlights substantial reward differences between preference pairs, which is crucial for effective preference learning. Collectively, these statistics demonstrate the dataset’s quality and its suitability for training robust process-level decision policies in agentic RAG frameworks.

2.3 Process-Supervised Preference Optimization

Based on the process-supervised preference data, we apply DPO [31] to tune the policy model. The optimization objective can be denoted as follows:

$$\mathcal{L}(\theta) = -\mathbb{E}_{(x, y_{< t}, y_t^w, y_t^l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{p_\theta(y_t^w | x, y_{< t})}{p_\theta(y_t^l | x, y_{< t})} \right) \right] \quad (3)$$

where x denotes the original question, $y_{< i}$ denote the the responses from previous reasoning steps. y_i^w and y_i^l represent the preferred and dispreferred responses in the subsequent step, respectively, and the hyperparameter β controls the KL constraint.

2.4 Agentic RAG Inference

To enable LLMs to autonomously interact with external information, we propose an agentic RAG workflow that supports adaptive reasoning through iterative search and reflection. ReasonRAG allows the model to dynamically determine when and how to invoke search engine based on question complexity.

The workflow operates through three recursive decision states: Reasoning, Grounding, and Terminal. In the **Reasoning** state, the LLM evaluates the current context to decide if it has sufficient information to answer the question. If sufficient, it generates a final answer enclosed in placeholders (<answer> and </answer>), thus terminating the process. If not, the model creates a new query enclosed in <query> and </query> placeholders to retrieve additional evidence. The system then transitions to the **Grounding** state, where documents are retrieved based on the query, and the model extracts relevant evidence spans. These evidence spans are appended to the context, after which the process loops back to the **Reasoning** state for further deliberation. (See Appendix G for detailed prompt designs.)

In summary, ReasonRAG supports multi-step, flexible reasoning while maintaining structured decision control. The use of explicit placeholders enhances interpretability and facilitates programmatic control during deployment. The complete algorithmic flow is provided in Algorithm 1.

Algorithm 1 Agentic RAG Inference Pipeline

Require: Original question x , large language model π_θ , retriever $\mathcal{R}$, maximum reasoning round N

Ensure: Final response y .

```

1: Initialize reasoning count  $i \leftarrow 0$ , and
   stage  $\leftarrow$  Reasoning
2: while  $i < N$  do
3:   if stage is Reasoning then
4:      $y_i \sim \pi_\theta(\cdot | x, y_{< i}, p_{\text{stage}})$ 
5:   else
6:      $y_i \sim \pi_\theta(\cdot | x, y_{< i}, docs, p_{\text{stage}})$ 
7:   end if
8:    $u \leftarrow u + u_i$ 
9:   if <query> detected in  $y_i$  then
10:    stage  $\leftarrow$  Grounding
11:     $q \leftarrow \text{extract\_query}(y_i)$ 
12:    docs  $\leftarrow \mathcal{R}(q)$ 
13:   else if <answer> detected in  $y_i$  then
14:     stage  $\leftarrow$  Terminal
15:     return extract_answer( $u_i$ )
16:   else if <evidence> detected in  $y_i$  then
17:     stage  $\leftarrow$  Reasoning
18:   end if
19:    $i \leftarrow i + 1$ 
20: end while
21: return final response  $y$ .
```

3 Experiments

3.1 Experimental Setup

Evaluation Dataset & Metrics. We evaluate ReasonRAG and all baselines on five public benchmarks: the single-hop QA dataset PopQA [32] and four multi-hop QA datasets, including HotpotQA [33], 2WikiMultiHopQA [34], Bamboogle [36], and MuSiQue [37]. Bamboogle and MuSiQue serve as out-of-domain QA evaluation datasets. The diversity of these datasets enables a comprehensive assessment of agentic RAG. We report Exact Match (EM) and F_1 scores as evaluation metrics. Refer to Appendix E for more details about dataset introduction, statistics, and metrics.

Baselines. We implement 12 baseline models which can be categorized into 6 types as follows: **Zero-shot:** Directly use prompt engineering on LLM to answer the question without or with retrieved documents [38]. **Active:** Actively make additional retrieval when retrieved data or generated responses have low confidence [39, 40]. **Adaptive:** Dynamically chooses the most suitable RAG pipeline from no-retrieval, single-hop, or multi-hop retrieval strategies [41]. **RAG-CoT:** Integrates

Type	Method	PopQA		HotpotQA		2WikiMulti		Bamboogle		MuSiQue		Avg.	
		EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁
Zero-shot	Naïve Generation	12.7	16.5	15.7	24.8	20.2	28.0	6.4	17.4	2.7	10.2	11.5	19.4
	Standard RAG	38.4	44.7	29.3	39.9	29.4	36.3	17.6	24.1	6.7	15.1	24.3	32.0
Active	FLARE	14.3	17.6	18.1	25.7	27.9	32.8	12.0	20.8	4.3	12.6	15.3	21.9
	Self-RAG(146k)	22.7	33.9	21.0	29.7	12.0	25.2	1.6	10.9	4.6	13.3	12.4	22.6
Adaptive	AdaptiveRAG(3k)	36.6	41.5	29.1	40.7	24.2	33.4	18.4	26.1	6.9	14.3	23.0	31.2
RAG-CoT	Iter-Retgen	38.7	44.9	30.3	42.1	31.2	38.7	19.2	26.4	7.7	14.2	25.4	33.3
	IRCoT	36.2	43.6	27.7	41.5	23.5	32.5	17.2	22.5	8.6	13.2	22.6	30.7
Summary	RECOMP	40.5	45.8	29.7	41.2	33.2	39.4	21.7	28.6	9.2	15.8	26.9	34.2
	LongLLMLingua	39.2	45.1	31.4	43.2	34.5	40.2	20.3	27.4	8.7	14.9	26.8	34.2
	Selective-Context	34.9	41.5	19.3	27.3	20.3	29.7	15.3	22.6	6.1	13.7	19.2	27.0
Reasoning	AutoRAG(10k)	38.6	44.1	33.3	43.7	39.5	46.1	24.8	32.2	11.3	18.3	29.5	36.9
	Search-R1(90k)	39.7	44.8	37.0	47.0	41.4	48.0	32.0	43.8	14.6	19.9	32.8	40.7
	ReasonRAG(5k)	41.5*	46.2*	38.4*	48.9*	43.6*	50.4*	36.0*	45.5*	12.8	20.6*	34.4*	42.3*

Table 2: Main Results (%) on Five benchmarks (the number of queries used for training is indicated in brackets). “*” indicates the statistical significance (i.e., two-sided t-test with $p < 0.05$) over the best baseline. Two most important columns: the averaged EM and F₁ are highlighted

chain-of-thought reasoning with retrieval, enabling multi-step, evidence-seeking answers [42, 43]. **Summary:** Compresses or summarizes retrieved content to fit model input constraints while retaining key information [44, 45, 46]. **Reasoning:** Enhances multi-hop reasoning by structuring the reasoning process and scrutinizing retrieved evidence [47, 25]. Note that ReasonRAG and all baselines use Qwen2.5-7B-Instruct [48] as the backbone model, ensuring fair comparison. Refer to more implementation details about ReasonRAG and baselines in Appendix F.1 and F.2.

3.2 Main Results

We present detailed performance results on ReasonRAG against 12 baselines across five benchmark datasets, as shown in Table 2. Our key findings are summarized below:

- **Data Efficiency:** ReasonRAG, despite being trained on only 5k queries, outperforms the search-R1 baseline trained with 90k queries. On average across all datasets, ReasonRAG achieves higher EM (34.4%) and F₁ (42.3%) scores compared to search-R1 (32.8% EM, 40.7% F₁), highlighting the superior data efficiency of ReasonRAG. This demonstrates the effectiveness of process-supervised RL, which leverages fine-grained rewards, over current outcome-supervised methods.
- **Multi-hop Reasoning:** ReasonRAG shows substantial performance gains on multi-hop reasoning tasks. On the HotpotQA dataset, it achieves an F₁ score of 48.9%, outperforming models like AutoRAG (43.7% F₁) and search-R1 (47.0% F₁), both of which are trained on larger datasets. This underscores ReasonRAG’s strength in handling complex, multi-step questions that require integrating evidence from multiple sources.
- **Out-of-domain Generalization:** ReasonRAG demonstrates strong generalization to out-of-domain data. On challenging benchmarks such as Bamboogle and MuSiQue, it consistently achieves higher F₁ scores relative to other baselines. This indicates improved robustness and transferability of its reasoning capabilities across different domains.

3.3 Training Efficiency

Figure 4 compares the training efficiency of ReasonRAG and Search-R1, both trained on 4 NVIDIA A100 80GB GPUs. The figure illustrates the progression of EM scores with increasing GPU hours across three datasets. The results reveal that ReasonRAG has higher training efficiency compared to Search-R1. ReasonRAG achieves superior EM scores with fewer GPU hours, indicating that it requires less training data and compute to reach strong performance levels. In contrast, Search-R1 requires significantly more GPU hours to reach similar performance.




Figure 4: Training cost and convergence speed comparison (EM%) for ReasonRAG and Search-R1

Method	PopQA		HotpotQA		2WikiMulti		Bamboogle		MuSiQue		Avg.	
	EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁	EM	F ₁
ReasonRAG (Base)	35.6	42.7	23.7	38.2	15.2	28.9	28.0	38.7	7.7	15.4	22.0	32.8
ReasonRAG (SFT)	31.6	37.4	26.8	38.7	35.1	40.9	17.6	27.3	8.6	15.5	23.9	32.0
ReasonRAG (RL-ORL): 5k queries	23.0	30.9	28.1	32.6	32.0	43.8	17.5	24.1	5.9	13.1	21.3	28.9
ReasonRAG (RL-ORL): 10k queries	39.5	45.7	36.7	46.7	40.5	47.2	30.7	40.6	12.6	19.5	32.0	39.9
ReasonRAG (RL-PRL)	41.5	46.2	38.4	48.9	43.6	50.4	36.0	45.5	12.8	20.6	34.5	42.3

Table 3: Impact of different optimization strategies on ReasonRAG’s effectiveness.

The efficiency gap between the two models is particularly notable on multi-hop question answering tasks. For the single-hop PopQA dataset, performance gains for both models are comparably rapid as training progresses. However, for multi-hop datasets such as 2WikiMultiHopQA and HotpotQA, ReasonRAG consistently demonstrates significant improvements with increased GPU hours. This further underscores its effectiveness on complex reasoning tasks, where it delivers faster and greater performance improvements with fewer resources.

3.4 Effectiveness of Different Optimization Strategies

In this section, we compare the effectiveness of ReasonRAG utilizing three different optimization strategies against the base model. Our default approach, ReasonRAG (RL-PRL), is trained with process-level rewards as described in Section 2. For ReasonRAG (RL-ORL), we adopt outcome-level reward training following the Search-R1 protocol [25]. Specifically, we evaluate two versions: RL-ORL-5k, trained on the same 5k queries as RL-PRL, and RL-ORL-10k, which incorporates an additional 5k queries sampled from PopQA, HotpotQA, and 2WikiMultiHopQA, totaling 10k training examples. For ReasonRAG (SFT), we use the preferred responses from the RAG-ProGuide preference pairs as ground truth and apply supervised fine-tuning (SFT) via next-token prediction. Table 3 summarizes the performance of these four variants. Our main findings are as follows:

- **Superiority of PRL:** ReasonRAG (PRL) consistently outperforms all other variants across all datasets, both in-domain and out-of-domain, indicating stronger generalization capabilities.
- **High Data Demand of ORL:** ReasonRAG (ORL) achieves the second-best results, but requires substantially more training data to match the comparable performance of PRL. Although ORL is more effective than Base and SFT, its training efficiency is relatively low.
- **Overfitting in SFT:** SFT leads to overfitting on multi-hop reasoning paths, resulting in reduced performance on single-hop tasks. Furthermore, SFT-trained models generalize poorly, as demonstrated by a marked performance decline on the Bamboogle dataset.

3.5 Impact of Search on Performance

Performance vs. Retrieval Steps. Figure 5 shows the EM performance of ReasonRAG across varying retrieval iterations on 3 datasets. We observe a consistent trend: performance improves with more retrieval iterations and then gradually saturates. Notably, ReasonRAG can adaptively determine the required inference depth according to task complexity. For the single-hop PopQA dataset, performance converges within 2 to 3 retrieval steps, whereas more complex multi-hop tasks such as 2WikiMultiHopQA and HotpotQA require 3 to 5 steps to reach peak performance. In contrast, ReasonRAG (base) without preference optimization only achieves reliable gains on PopQA and




Figure 5: The EM performance across varying retrieval iterations on 3 benchmarks.




Figure 6: Effect of top- k retrieved documents on ReasonRAG’s performance across 3 datasets.

struggles to handle multi-hop settings. These results demonstrate ReasonRAG’s ability to perform adaptive reasoning in response to the complexity of the input question.

Performance vs. Number of Retrieved Documents. Figure 6 compares the performance of ReasonRAG under different top- k retrieval settings, where k refers to retrieving top- k relevant document passages per search query. Results indicate that while ReasonRAG remains robust across a range of k values, its performance is sensitive to the quantity of retrieved information. With $k=1$, limited context restricts the model’s reasoning ability. Increasing k to 3 yields significant improvements across all datasets, suggesting that ReasonRAG effectively leverages additional evidence. Further increasing k to 5 does not further increase the performance on PopQA and HotpotQA, whereas 2WikiMultiHopQA continues to benefit from richer retrieved context. Overall, these findings highlight ReasonRAG’s capacity to utilize additional retrieved documents, particularly in more complex multi-hop scenarios.

4 Related Works

Prompt-Based Agentic RAG. Early prompt-based approaches leverage manually designed workflows to elicit the inherent capabilities of LLMs for interacting with external information. Specifically, the RAG task is often decomposed into subtasks such as adaptive retrieval judgment [41], query generation [19, 42, 43], evidence extraction [40, 44, 46, 45], and answer generation [39]. Recently, agentic RAG aims to design workflows that empower LLM to autonomously interact with external information. OpenResearcher [49], AirRAG [50], IterDRAG [51], PlanRAG [52], and search-ol [53] demonstrate strong performance improvement by the effective incorporation with the search engine. Nevertheless, these methods are limited by their dependence on inherent capabilities for interacting with external information and the requirement for manual design when applied to new domains.

RL-Based Agentic RAG. Outcome-supervised RL has rapidly emerged as a mainstream paradigm for enhancing LLM reasoning capabilities by providing end-to-end optimization. Following the widespread adoption of RL by major AI companies to improve the reasoning abilities of their models on complex tasks [54, 55, 56], recent work [25, 57, 58, 59] has extended outcome-supervised reinforcement learning to RAG, empowering LLMs to autonomously utilize search engines for intricate inference. While outcome-supervised RL has demonstrated significant performance gains, it also faces challenges such as reward sparsity, training instability, and substantial computational cost. Moreover, outcome-supervised RL typically demands extensive training resources. In contrast, process-supervised RL has recently been applied to enhance reasoning abilities, outperforming

outcome-supervised approaches by providing fine-grained rewards [60, 61, 62, 63]. As an alternative avenue for improving LLM reasoning in RAG, process-supervised RL for RAG remains unexplored.

5 Conclusion

We introduce ReasonRAG, a process-supervised agentic RAG method for fine-grained policy optimization. Our approach integrates Monte Carlo Tree Search (MCTS) with the agentic RAG workflow to generate RAG-ProGuide, a high-quality dataset providing process-level supervision by prioritizing the shortest reasoning paths leading to correct answers. Leveraging RAG-ProGuide, we perform preference-based policy optimization to enhance LLMs’ autonomous capabilities in query generation, evidence extraction, and answer synthesis. Experiments demonstrate that ReasonRAG achieves superior performance on five benchmark datasets using only 5k training instances, significantly fewer than the 90k required by Search-R1, highlighting the effectiveness of RAG-ProGuide’s high-quality process-level rewards in optimizing agentic RAG policies.

References

- [1] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- [2] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. Palm: Scaling language modeling with pathways. *Journal of Machine Learning Research*, 24(240):1–113, 2023.
- [3] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [4] Yizhang Zhu, Shiyin Du, Boyan Li, Yuyu Luo, and Nan Tang. Are large language models good statisticians? In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.
- [5] Javier González and Aditya Nori. Does reasoning emerge? examining the probabilities of causation in large language models. *Advances in Neural Information Processing Systems*, 37:117737–117761, 2024.
- [6] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM computing surveys*, 55(12):1–38, 2023.
- [7] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, et al. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Transactions on Information Systems*, 43(2):1–55, 2025.
- [8] Nick McKenna, Tianyi Li, Liang Cheng, Mohammad Hosseini, Mark Johnson, and Mark Steedman. Sources of hallucination by large language models on inference tasks. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- [9] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks, 2021.
- [10] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Haofen Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2, 2023.

- [11] Wenqi Fan, Yujuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin, Tat-Seng Chua, and Qing Li. A survey on rag meeting llms: Towards retrieval-augmented large language models. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 6491–6501, 2024.
- [12] Guangzhi Xiong, Qiao Jin, Zhiyong Lu, and Aidong Zhang. Benchmarking retrieval-augmented generation for medicine. In *Findings of the Association for Computational Linguistics ACL 2024*, pages 6233–6251, 2024.
- [13] Aditi Singh, Abul Ehtesham, Saket Kumar, and Tala Talaei Khoei. Agentic retrieval-augmented generation: A survey on agentic rag. *arXiv preprint arXiv:2501.09136*, 2025.
- [14] Google. Gemini deep research. 2024.
- [15] OpenAI. Deep research system card. *Technical Report*, 2025.
- [16] Maojia Song, Shang Hong Sim, Rishabh Bhardwaj, Hai Leong Chieu, Navonil Majumder, and Soujanya Poria. Measuring and enhancing trustworthiness of llms in rag through grounded attributions and learning to refuse. In *The Thirteenth International Conference on Learning Representations*.
- [17] Jaehyung Kim, Jaehyun Nam, Sangwoo Mo, Jongjin Park, Sang-Woo Lee, Minjoon Seo, Jung-Woo Ha, and Jinwoo Shin. Sure: Summarizing retrievals using answer candidates for open-domain qa of llms. In *The Twelfth International Conference on Learning Representations*.
- [18] Zhiqing Sun, Xuezhi Wang, Yi Tay, Yiming Yang, and Denny Zhou. Recitation-augmented language models. In *The Eleventh International Conference on Learning Representations*.
- [19] Reinald Kim Amplayo, Kellie Webster, Michael Collins, Dipanjan Das, and Shashi Narayan. Query refinement prompts for closed-book long-form qa. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7997–8012, 2023.
- [20] Yuhao Wang, Ruiyang Ren, Junyi Li, Wayne Xin Zhao, Jing Liu, and Ji-Rong Wen. Rear: A relevance-aware retrieval-augmented framework for open-domain question answering. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 5613–5626, 2024.
- [21] Biao Zhang, Zhongtao Liu, Colin Cherry, and Orhan Firat. When scaling meets llm finetuning: The effect of data, model and finetuning method. In *The Twelfth International Conference on Learning Representations*.
- [22] Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang Tong, Saining Xie, Dale Schuurmans, Quoc V Le, Sergey Levine, and Yi Ma. Sft memorizes, rl generalizes: A comparative study of foundation model post-training. *arXiv preprint arXiv:2501.17161*, 2025.
- [23] Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, et al. Openai o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.
- [24] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-rl: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- [25] Bowen Jin, Hansi Zeng, Zhenrui Yue, Jinsung Yoon, Sercan Arik, Dong Wang, Hamed Zamani, and Jiawei Han. Search-rl: Training llms to reason and leverage search engines with reinforcement learning. *arXiv preprint arXiv:2503.09516*, 2025.
- [26] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *The Twelfth International Conference on Learning Representations*, 2023.
- [27] Yuqing Du, Alexander Havrilla, Sainbayar Sukhbaatar, Pieter Abbeel, and Roberta Raileanu. A study on improving reasoning in language models. In *I Can’t Believe It’s Not Better Workshop: Failure Modes in the Age of Foundation Models*, 2024.

- [28] Liangchen Luo, Yinxiao Liu, Rosanne Liu, Samrat Phatale, Meiqi Guo, Harsh Lara, Yunxuan Li, Lei Shu, Yun Zhu, Lei Meng, et al. Improve mathematical reasoning in language models by automated process supervision. *arXiv preprint arXiv:2406.06592*, 2024.
- [29] Jeremy Scheurer, Jon Ander Campos, Jun Shern Chan, Angelica Chen, Kyunghyun Cho, and Ethan Perez. Learning from natural language feedback. In *ACL Workshop on Learning with Natural Language Supervision*, 2022.
- [30] Cameron B Browne, Edward Powley, Daniel Whitehouse, Simon M Lucas, Peter I Cowling, Philipp Rohlfshagen, Stephen Tavener, Diego Perez, Spyridon Samothrakis, and Simon Colton. A survey of monte carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in games*, 4(1):1–43, 2012.
- [31] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36:53728–53741, 2023.
- [32] Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Hannaneh Hajishirzi, and Daniel Khashabi. When not to trust language models: Investigating effectiveness and limitations of parametric and non-parametric memories. *arXiv preprint arXiv:2212.10511*, 7, 2022.
- [33] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D Manning. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2018.
- [34] Xanh Ho, Anh-Khoa Duong Nguyen, Saku Sugawara, and Akiko Aizawa. Constructing a multi-hop qa dataset for comprehensive evaluation of reasoning steps. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 6609–6625, 2020.
- [35] Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.
- [36] Ofir Press, Muru Zhang, Sewon Min, Ludwig Schmidt, Noah A Smith, and Mike Lewis. Measuring and narrowing the compositionality gap in language models. In *The 2023 Conference on Empirical Methods in Natural Language Processing*.
- [37] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. Musique: Multihop questions via single-hop question composition. *Transactions of the Association for Computational Linguistics*, 10:539–554, 2022.
- [38] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems*, 33:9459–9474, 2020.
- [39] Zhengbao Jiang, Frank F Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. Active retrieval augmented generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7969–7992, 2023.
- [40] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-rag: Learning to retrieve, generate, and critique through self-reflection. In *The Twelfth International Conference on Learning Representations*.
- [41] Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, and Jong-Cheol Park. Adaptive-rag: Learning to adapt retrieval-augmented large language models through question complexity. In *2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 7036–7050. Association for Computational Linguistics, 2024.

- [42] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *The 61st Annual Meeting Of The Association For Computational Linguistics*, 2023.
- [43] Zhihong Shao, Yeyun Gong, Minlie Huang, Nan Duan, Weizhu Chen, et al. Enhancing retrieval-augmented large language models with iterative retrieval-generation synergy. In *The 2023 Conference on Empirical Methods in Natural Language Processing*.
- [44] Fangyuan Xu, Weijia Shi, and Eunsol Choi. Recomp: Improving retrieval-augmented lms with compression and selective augmentation. In *12th International Conference on Learning Representations, ICLR 2024*, 2024.
- [45] Huiqiang Jiang, Qianhui Wu, Xufang Luo, Dongsheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. Longllmlingua: Accelerating and enhancing llms in long context scenarios via prompt compression. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1658–1677, 2024.
- [46] Yucheng Li, Bo Dong, Frank Guerin, and Chenghua Lin. Compressing context to enhance inference efficiency of large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6342–6353, 2023.
- [47] Tian Yu, Shaolei Zhang, and Yang Feng. Auto-rag: Autonomous retrieval-augmented generation for large language models. *arXiv preprint arXiv:2411.19443*, 2024.
- [48] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- [49] Yuxiang Zheng, Shichao Sun, Lin Qiu, Dongyu Ru, Cheng Jiayang, Xuefeng Li, Jifan Lin, Binjie Wang, Yun Luo, Renjie Pan, et al. Openresearcher: Unleashing ai for accelerated scientific research. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 209–218, 2024.
- [50] Wenfeng Feng, Chuzhan Hao, Yuewei Zhang, Jingyi Song, and Hao Wang. Airrag: Activating intrinsic reasoning for retrieval augmented generation via tree-based search. *arXiv preprint arXiv:2501.10053*, 2025.
- [51] Zhenrui Yue, Honglei Zhuang, Aijun Bai, Kai Hui, Rolf Jagerman, Hansi Zeng, Zhen Qin, Dong Wang, Xuanhui Wang, and Michael Bendersky. Inference scaling for long-context retrieval augmented generation. *arXiv preprint arXiv:2410.04343*, 2024.
- [52] Prakhar Verma, Sukruta Prakash Midigeshi, Gaurav Sinha, Arno Solin, Nagarajan Natarajan, and Amit Sharma. Plan rag: Efficient test-time planning for retrieval augmented generation. In *Workshop on Reasoning and Planning for Large Language Models*, 2025.
- [53] Xiaoxi Li, Guanting Dong, Jiajie Jin, Yuyao Zhang, Yujia Zhou, Yutao Zhu, Peitian Zhang, and Zhicheng Dou. Search-o1: Agentic search-enhanced large reasoning models. *arXiv preprint arXiv:2501.05366*, 2025.
- [54] Kimi Team, Angang Du, Bofei Gao, Bowei Xing, Changjiu Jiang, Cheng Chen, Cheng Li, Chenjun Xiao, Chenzhuang Du, Chonghua Liao, et al. Kimi k1. 5: Scaling reinforcement learning with llms. *arXiv preprint arXiv:2501.12599*, 2025.
- [55] DiJia Su, Sainbayar Sukhbaatar, Michael Rabbat, Yuandong Tian, and Qingqing Zheng. Dual-former: Controllable fast and slow thinking by learning with randomized reasoning traces. In *The Thirteenth International Conference on Learning Representations*, 2024.
- [56] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Y Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- [57] Yuxiang Zheng, Dayuan Fu, Xiangkun Hu, Xiaojie Cai, Lyumanshan Ye, Pengrui Lu, and Pengfei Liu. Deepresearcher: Scaling deep research via reinforcement learning in real-world environments. *arXiv preprint arXiv:2504.03160*, 2025.

- [58] Hongru Wang, Cheng Qian, Wanjun Zhong, Xiushi Chen, Jiahao Qiu, Shijue Huang, Bowen Jin, Mengdi Wang, Kam-Fai Wong, and Heng Ji. Otc: Optimal tool calls via reinforcement learning. *arXiv preprint arXiv:2504.14870*, 2025.
- [59] Hao Sun, Zile Qiao, Jiayan Guo, Xuanbo Fan, Yingyan Hou, Yong Jiang, Pengjun Xie, Fei Huang, and Yan Zhang. Zerossearch: Incentivize the search capability of llms without searching. *arXiv preprint arXiv:2505.04588*, 2025.
- [60] Ruotian Ma, Peisong Wang, Cheng Liu, Xingyan Liu, Jiaqi Chen, Bang Zhang, Xin Zhou, Nan Du, and Jia Li. S2r: Teaching llms to self-verify and self-correct via reinforcement learning. *arXiv preprint arXiv:2502.12853*, 2025.
- [61] Zongyu Lin, Yao Tang, Xingcheng Yao, Da Yin, Ziniu Hu, Yizhou Sun, and Kai-Wei Chang. Qlass: Boosting language agent inference via q-guided stepwise search. *arXiv preprint arXiv:2502.02584*, 2025.
- [62] Peiyi Wang, Lei Li, Zhihong Shao, Runxin Xu, Damai Dai, Yifei Li, Deli Chen, Yu Wu, and Zhifang Sui. Math-shepherd: Verify and reinforce llms step-by-step without human annotations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9426–9439, 2024.
- [63] Muhammad Khalifa, Rishabh Agarwal, Lajanugen Logeswaran, Jaekyeom Kim, Hao Peng, Moontae Lee, Honglak Lee, and Lu Wang. Process reward models that think. *arXiv preprint arXiv:2504.16828*, 2025.
- [64] Levente Kocsis and Csaba Szepesvári. Bandit based monte-carlo planning. In *European conference on machine learning*, pages 282–293. Springer, 2006.
- [65] Jiajie Jin, Yutao Zhu, Xinyu Yang, Chenghao Zhang, and Zhicheng Dou. Flashrag: A modular toolkit for efficient retrieval-augmented generation research. *CoRR*, abs/2405.13576, 2024.

Appendix Overview

The appendix includes the following sections:

- **Appendix A:** Details the MCTS tree construction process using process-level rewards. This serves as an extended explanation to Section 2.2.2.
- **Appendix B:** Illustrates the example of LLM evaluation for process-level simulation. This serves as an extended example of MCTS node expansion in Section 2.2.2.
- **Appendix C:** Compares the reasoning response between ReasonRAG(base) and ReasonRAG(PRL). This serves as an extended example for Section 3.4.
- **Appendix D:** Visualizes our generated process-level reward annotations. This serves as an extended example for Section 2.2.3.
- **Appendix E:** Provides additional details of the evaluation setup. This serves as a supplement to Section 3.1.
- **Appendix F:** Provide more details of the implementation. This serves as a supplement to Section 3.1.
- **Appendix G:** Details the prompt design for Agentic RAG. This serves a supplement to Section 2.4.
- **Appendix H:** Details the licensing terms and conditions governing the use and distribution of the proposed datasets.
- **Appendix I:** Discusses the limitations and constraints of the proposed approach.
- **Appendix J:** Evaluates the potential societal implications and ethical considerations of the research.

A Monte Carlo Tree Search for Process-level Reward

Formally, for each RAG intermediate process, its corresponding state s encompasses the original question x , the preceding thoughts $y_{<i}$, and a stage $\in \{\text{Reasoning, Grounding, Terminal}\}$. The stage indicates the current decision mode within the reasoning process. In the Reasoning stage, the LLM autonomously decides whether to generate a new query or directly produce an answer. If a query is chosen, it triggers a call to the external search engine, and the retrieved documents are added to the context in the next step. If the model opts to generate an answer instead, the process transitions into the Terminal stage. In the Grounding stage, the model extracts relevant evidence spans from the retrieved documents based on the most recent query. After extracting evidence, the state transitions back to the Reasoning stage, enabling further iterative reasoning.

Based on the state, the policy of the next action is defined as

$$\pi(a|s) = \text{LLM}(a|s) = \begin{cases} \pi_\theta(\cdot|x, y_{<i}, p_{\text{stage}}), & \text{if stage is Reasoning} \\ \pi_\theta(\cdot|x, y_{<i}, \text{docs}, p_{\text{stage}}), & \text{otherwise} \end{cases} \quad (4)$$

Consequently, the state transition can be represented as $s_{t+1} = \text{concatenate}(s_t, a_t)$. Each node in the tree contains the following information: $\{N(s), Q(s), \text{Stage}(s)\}$, where $N(s)$ denotes the number of times state s has been visited. $Q(s)$ represents the current intermediate annotation collected through the Monte Carlo method, with values in the range $[0, 1]$. With the tree structure defined, the MCTS begins from the root node and constructs the tree by iteratively performing three key operations: selection, expansion, and backpropagation.

Selection: This step aims to select nodes that balance the search quality and exploration degree. The node selection starts from the root node, and iteratively selects the child nodes based on their state value Q and visiting frequency N . These variable are refined during the search strategy, detailed in the backpropagation section. To effectively trade off between exploring unvisited nodes and exploiting nodes with higher state value, we iteratively search for the next node using UCT score [64]. The state is selected according to the following formula:

$$s_i^* = \operatorname{argmax}_{s_i \in \text{children}(s)} [Q(s_i) + c_{uct} \frac{\sqrt{\sum_i N(s_i)}}{1 + N(s_i)}] \quad (5)$$

where c_{uct} is a trade-off parameter to control the exploration degree. The algorithm begins by exploring unvisited states and progressively favors nodes with higher Q -values and fewer visits.

Expansion: Given a selected node that does not reach a terminal state and maximum child node limit, the expansion step proceeds with a single step of RAG reasoning based on Equation (4) and initializes a new child node with the generated response. Following response generation, the simulation step iteratively reasons until a final answer is reached, serving as the basis for initializing the reward of the created node. However, simulating RAG leads to inefficiency due to its need for iterative LLM reasoning and retrieval. To address this challenge, the correctness of the intermediate reasoning process is evaluated by LLM judgment based on the intermediate process against the golden answer, outputting a correctness value $v \in [0, 1]$, as defined in Equation (6). This approach avoids time-consuming simulations, providing an efficient evaluation for exploring the complex RAG reasoning space.

$$v(s) = \pi_{\theta}(\cdot | x, y_{1:t}, \text{answer}, p_{\text{eval}}) \quad (6)$$

Backpropagation: Following node expansion, the node statistics are iteratively updated from the current node back to the root node. The correctness score $Q(s_t)$ is updated based on Eq. (1). The visit count for the traversed nodes is incremented:

$$N(s) \leftarrow N(s) + 1 \quad \forall s \in \text{path from } s_t \text{ to root} \quad (7)$$

B Process Evaluation

The three cases illustrated in Figures 7–9 demonstrate the effectiveness of ReasonRAG in supporting stepwise reasoning over retrieved content. In the partially correct case, the agent retrieves relevant evidence but fails to cover all required entities, leading to incomplete inference. In the incorrect case, the agent misidentifies directors, resulting in a reasoning chain based on factual errors. In contrast, the fully correct case shows ReasonRAG’s ability to decompose the question into sequential sub-queries, accurately retrieve biographical information, and synthesize the correct answer through structured comparison. Although minor imperfections exist in the query formulation, the reasoning trace demonstrates clear progression, factual grounding, and answer justification. These cases collectively highlight ReasonRAG’s advantage in aligning sub-question formulation with retrieval intent and maintaining semantic coherence across reasoning steps, leading to improved robustness and interpretability in complex multi-hop scenarios.

Case of LLM judge partial correct process

Question
 “Which film has the director died later, Il Coraggio or Shark Monroe?”

Golden Answers
 “Il Coraggio.”

Agent Thoughts:
 So the next query is <query> What is the release year of Il Coraggio and who directed it? </query> Based on the query, the relevant evidence is <evidence>Il Coraggio is a 1955 Italian comedy film directed by Domenico Paolella. </evidence>

LLM judge

The agent’s reasoning is severely incomplete and structurally flawed. It begins with a partial query that only addresses Il Coraggio but completely ignores Shark Monroe, (Omitted for brevity)

Moreover, the query does not directly request the death dates of the directors, which is the core requirement (Omitted for brevity)

While the evidence retrieved is correct and related to Il Coraggio, it is insufficient (Omitted for brevity)

So the score is 15.

Figure 7: An example of agent partial correctness in a RAG process.

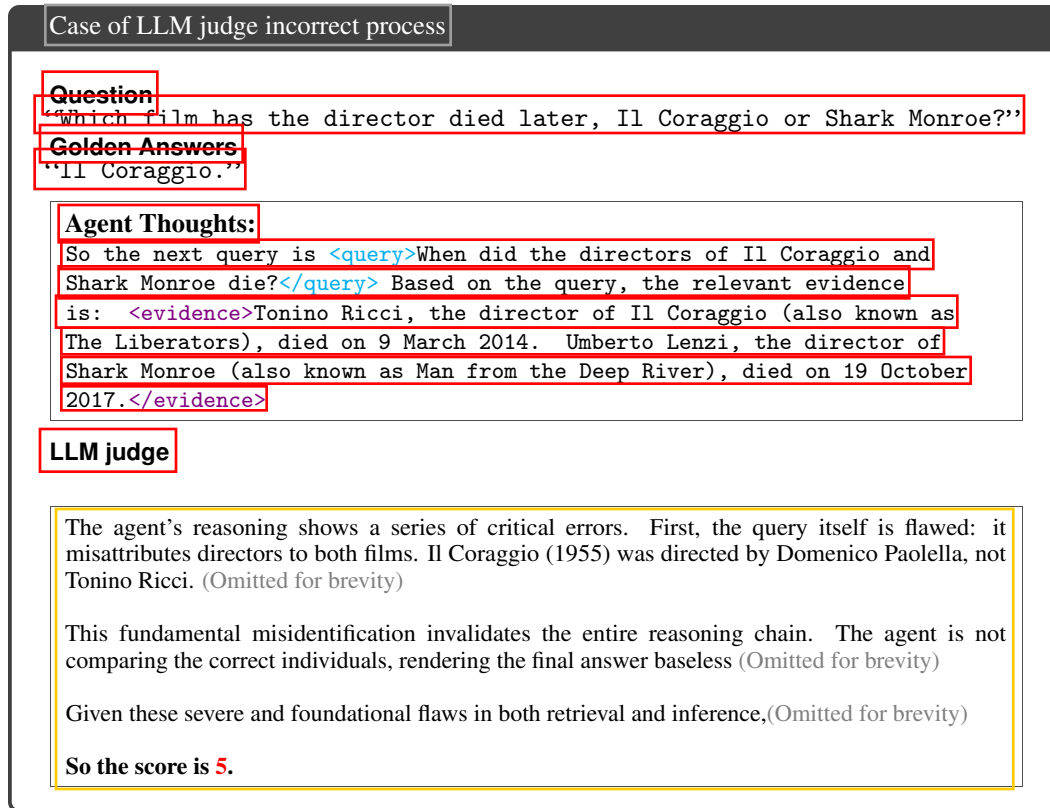


Figure 8: An example of agent incorrectness in a RAG process.

C Case Study

Figure 10 shows an example for agentic RAG response from ReasonRAG. Before the policy optimization, LLMs fails to generate the appropriate query and mislead by irrelevant information. In contrast, the process-supervised RL empower LLMs to autonomously invoke query generation, evidence extraction, and answer generation.

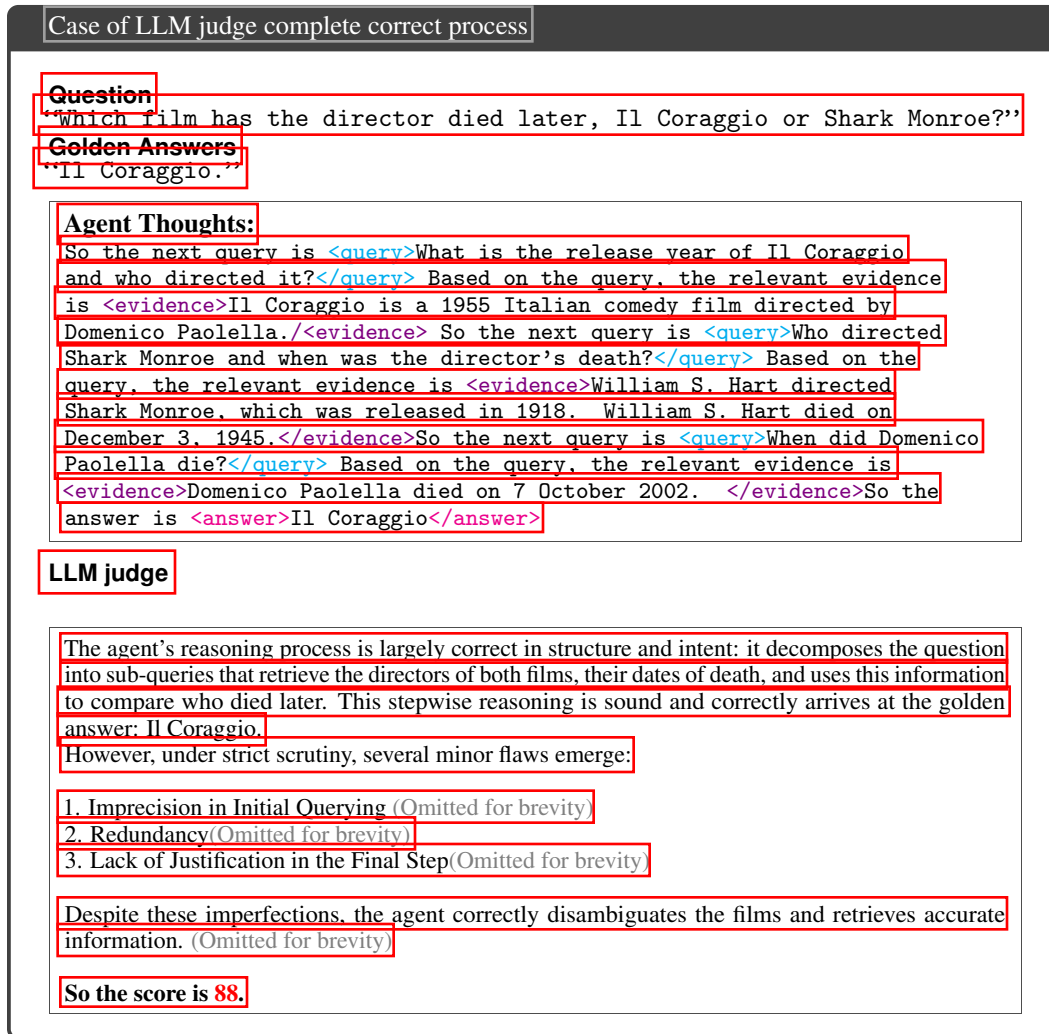


Figure 9: An example of agent complete correctness in a RAG process.

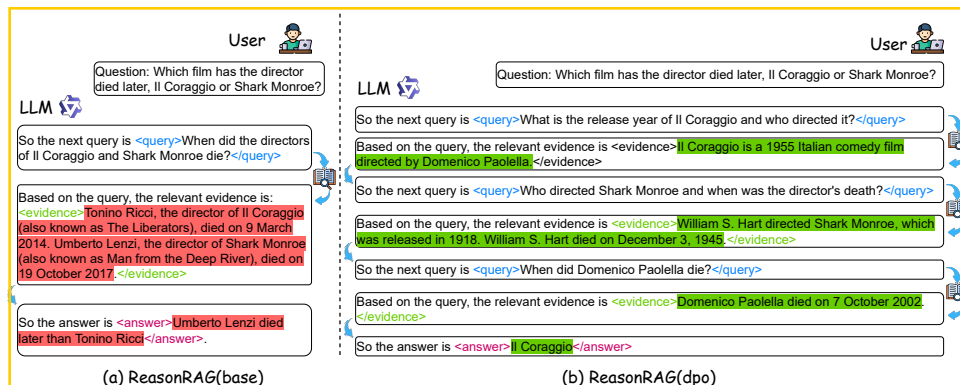


Figure 10: Case Study on 2wiki dataset.




Figure 11: A Tree Example for the process-level preference annotations.

D Example for constructed tree

Figure 11 illustrates an example for the tree-structured process data. The root node corresponding to the original question, each node corresponds one-step response from LLMs. The correct reasoning path has been colored in orange and annotated with a higher reward.

E Evaluation Dataset Details

Evaluation Datasets. We use the process-level annotated questions (4,603 examples) as the training set for our policy optimization. These include 704 questions from PopQA, 2,843 from HotpotQA, and 1,056 from 2WikiMultiHopQA. For evaluation, we use the remaining unlabeled samples from PopQA as the test set, and we adopt the official development splits of HotpotQA and 2WikiMultiHopQA as test sets for multi-hop reasoning evaluation. Table 4 summarizes the training and test set sizes for each source. These datasets vary in their design focus and reasoning requirements. HotpotQA and 2WikiMultiHopQA are constructed to evaluate multi-hop reasoning capabilities, where answering a question requires combining information from multiple passages. HotpotQA includes sentence-level supporting facts and diverse question types, such as bridge and comparison questions. 2WikiMultiHopQA ensures genuine multi-step inference by leveraging structured knowledge from Wikidata and constructing explicit reasoning paths. PopQA, in contrast, is an open-domain QA dataset designed to probe factual recall in large language models. It focuses on a wide spectrum of factual knowledge, from high-frequency popular facts to long-tail, less commonly known information. By combining these datasets, we cover a diverse set of reasoning challenges, including factual retrieval, multi-hop inference, and process-level supervision.

Dataset Source	Train Set Size	Test Set Size
PopQA	704	11,267
HotpotQA	2,843	7,405
2WikiMultiHopQA	1,056	12,576
bamboogle	-	125
musique	-	2,417
Total	4,603	33,790

Table 4: Number of examples in the training and test sets for each dataset. Process-level annotations are used for training; test sets include remaining PopQA examples and official development splits of other datasets.

Evaluation Details. To evaluate model performance on question answering tasks, we adopt two standard metrics: Exact Match (EM) and F_1 score.

Name	Purpose	Artifact URL
PopQA	Eval Dataset	https://huggingface.co/datasets/RUC-NLPIR/FlashRAG_datasets/tree/main/popqa
HotpotQA	Eval Dataset	https://huggingface.co/datasets/RUC-NLPIR/FlashRAG_datasets/tree/main/hotpotqa
2WikiMultiHopQA	Eval Dataset	https://huggingface.co/datasets/RUC-NLPIR/FlashRAG_datasets/tree/main/2wikimultihopqa
Bamboogle	Eval Dataset	https://huggingface.co/datasets/RUC-NLPIR/FlashRAG_datasets/tree/main/bamboogle
MuSiQue	Eval Dataset	https://huggingface.co/datasets/RUC-NLPIR/FlashRAG_datasets/tree/main/musique
RAG-ProGuide	Train Dataset	https://anonymous.4open.science/r/ReasonRAG-B442
BGE	Retriever	https://huggingface.co/BAAI/bge-base-en-v1.5
Wikidump 2018	Knowledge Source	https://archive.org/download/enwiki-20181220/enwiki-20181220-pages-articles.xml.bz2
Qwen2.5-7B-Instruct	Backbone Model	https://huggingface.co/Qwen/Qwen2.5-7B-Instruct
Adaptive-RAG	Baseline	https://huggingface.co/illuminoplanet/combined_flan_t5_xl_classifier
Self-RAG	Baseline	https://huggingface.co/selfrag/selfrag_llama2_13b
AutoRAG	Baseline	https://huggingface.co/ICTNLP/Auto-RAG-Llama-3-8B-Instruct
Search-R1	Baseline	https://huggingface.co/PeterJinGo/SearchR1-nq-hotpotqa_train-qwen2.5-7b-it-em-ppo
ReasonRAG	Our Method	https://anonymous.4open.science/r/ReasonRAG-B442

Table 5: Resource links of artifacts used in our experiments.

Exact Match (EM) measures the percentage of predictions that exactly match any of the reference answers. Formally,

$$\text{EM} = \frac{1}{N} \sum_{i=1}^N \delta(y_i^{\text{pred}}, y_i^{\text{gold}}), \quad (8)$$

where N is the number of examples, and $\delta(a, b) = 1$ if $a = b$, otherwise 0.

F₁ score computes the token-level overlap between the predicted answer and the ground-truth answer.

Let T_i^{pred} and T_i^{gold} denote the sets of tokens in the predicted and gold answers, respectively:

$$\text{Precision}_i = \frac{|T_i^{\text{pred}} \cap T_i^{\text{gold}}|}{|T_i^{\text{pred}}|}, \quad (9)$$

$$\text{Recall}_i = \frac{|T_i^{\text{pred}} \cap T_i^{\text{gold}}|}{|T_i^{\text{gold}}|}, \quad (10)$$

$$F_1 = \frac{1}{N} \sum_{i=1}^N \frac{2 \cdot \text{Precision}_i \cdot \text{Recall}_i}{\text{Precision}_i + \text{Recall}_i}. \quad (11)$$

We follow the official evaluation metrics implementation provided by the FlashRAG toolkit [65].

F Implementation Details

We summarize all artifacts (datasets, models, baselines, external knowledge base, etc) used in our experiments and their resource links in Table 5.

F.1 Implementation Details of ReasonRAG

Following the setup in the FlashRAG toolkit, we use Wikidump 2018 as our knowledge source. To ensure retrieval quality, we augment our corpus by incorporating relevant content from the PopQA, HotpotQA, and 2WikiMultiHopQA datasets. All datasets are available on Huggingface. Subsequently, we employ BGE as our retriever, consistently retrieving the top 3 documents. For all methods not requiring fine-tuning, we use Qwen2.5-7B-Instruct as our baseline model for inference.

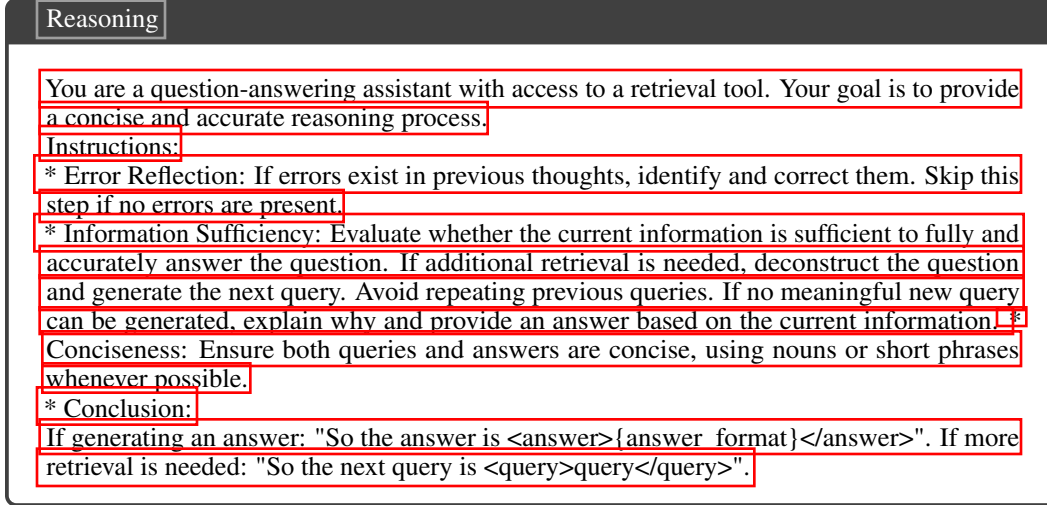


Figure 12: System prompt for Reasoning

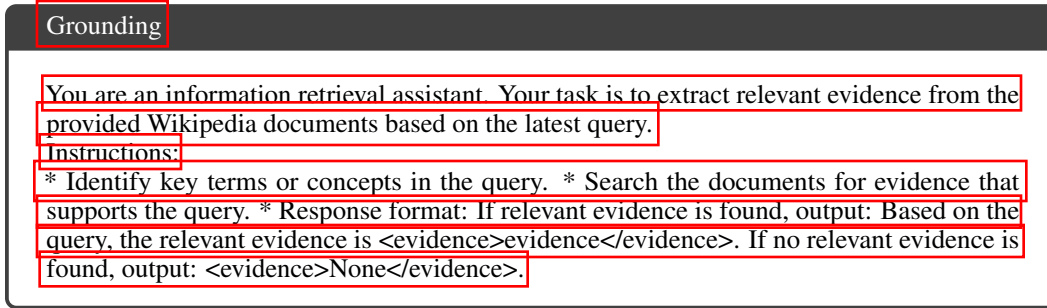


Figure 13: System prompt for Evidence Extraction

F.2 Implementation Details of Baselines

For baseline implementations, we utilize the FlashRAG [65] reproduction, where several models, such as Naïve Generation, Standard RAG, FLARE, Iter-Retgen, RECOMP, LongLLMLingua, and Selective-Context, require no additional parameter configuration. For Self-RAG, we use the checkpoint provided in the FlashRAG reproduction. For AdaptiveRAG, we employ the FlashRAG reproduction’s router and qwen2.5-7b-instruct as the reasoning model. For AutoRAG, we conduct inference using the publicly available checkpoint from Hugging Face. For Search-R1, we use the reproduced qwen2.5-7b-base and qwen2.5-7b-instruct checkpoints for inference.

G Prompt Instructions

Our prompts include Reasoning, Grounding in Agentic RAG workflow, and a process evaluation prompt for node expansion. No extra prompt design is needed when input a new question into LLMs for inference. The prompt details are illustrated in Figure 12, Figure 13, and Figure 14.

H License Agreement

The RAG-ProGuide is constructed based on popqa, hotpotqa, and 2wikimultihopqa from FlashRAG dataset [65]. All these datasets are using CC-BY-SA-4.0 license, allowing the modification for research use. For the new constructed dataset RAG-ProGuide, including but not limited to the questions and process-level reward, we make them available solely for research purposes. Users are permitted to use, modify, and share these annotations for academic and non-commercial research

Process Evaluation

An agent is tasked with answering a question using a retrieval tool. Critically assess its intermediate reasoning process to determine if it leads to the correct answer. Identify all flaws, inconsistencies, and mistakes in the thought process. Every imperfection, no matter how small, must be acknowledged. Evaluate how effectively the reasoning supports the final answer and the overall accuracy of the response. Ensure the evaluation is extremely harsh, leaving no leniency. Even if the answer seems close to correct, do not award full marks to maintain strict grading standards. Assign a score between $[0, 1]$ based on the severity of flaws and the reasoning's accuracy in leading to the golden answer. Respond briefly and conclude with: So the score is [Score].

Figure 14: System prompt for Process Evaluation

activities. Any other use, including commercial exploitation, is not permitted without explicit written permission from the authors.

I Limitations

We employ process-supervised RL to optimize the model policy. In contrast to outcome-supervised RL, our approach necessitates exploring process-level actions for fine-grained reward annotation. Consequently, acquiring sufficient data for process-level annotation incurs a higher time cost compared to outcome supervision during data rollout. Nevertheless, as the training efficiency verifies, our data exhibits superior quality, enabling models to achieve greater performance gains with fewer data samples.

J Societal Impacts

LLMs carry the risk of generating uncontrollable responses. When an LLM retrieves racist or harmful information from a search engine, it could be inadvertently led to produce similar content. We strongly advise users to employ our agentic RAG framework responsibly by integrating secure search engines or knowledge corpora and conducting evaluations within open-source, safe environments and datasets.